

The k-Nearest Neighbors (k-NN) Algorithm

The k-Nearest Neighbors (k-NN) algorithm is a simple, intuitive, and widely used method for both classification and regression tasks in machine learning.

Overview of k-Nearest Neighbors (k-NN)

How k-NN Works

Data Preparation: Gather and preprocess your dataset. Ensure features are normalized or scaled if they vary widely in range.

Choose k: Decide on the number of nearest neighbors (k) to consider.

Compute Distance: Calculate the distance between the query point (the point for which you want to make a prediction) and all points in the training set. Common distance metrics include **Euclidean**, Manhattan, and Minkowski distances.



Identify Neighbors: Select the k points in the training set that are closest to the query point based on the computed distance.

Make a Prediction: The query point is assigned to the class most common among its k nearest neighbors.


How to choose the value of k for KNN Algorithm?

The value of k is very crucial in the KNN algorithm to define the number of neighbors in the algorithm. The value of k in the k-nearest neighbors (k-NN) algorithm should be chosen based on the input data.

If the input data has more outliers or noise, a higher value of k would be better. It is recommended to choose an odd value for k to avoid ties in classification. Cross-validation methods can help in selecting the best k value for the given dataset.




0. Look at the data











Say you want to classify the grey point into a class. Here, there are three potential classes - lime green, green and orange.

1. Calculate distances









Start by calculating the distances between the grey point and all other points.

2. Find neighbours

Point Distance			
		2.1	→ 1st NN
		2.4	→ 2nd NN
		3.1	→ 3rd NN
		4.5	→ 4th NN

Next, find the nearest neighbours by ranking points by increasing distance. The nearest neighbours (NNs) of the grey point are the ones closest in dataspace.

3. Vote on labels

Class	# of votes	
	2	➔ Class  wins the vote! Point  is therefore predicted to be of class  .
	1	
	1	

Vote on the predicted class labels based on the classes of the k nearest neighbours. Here, the labels were predicted based on the $k=3$ nearest neighbours.

Euclidean Distance Equation

$$\text{Euclidean distance} = \sqrt{(f_{11} - f_{12})^2 + (f_{21} - f_{22})^2}$$

where f_{11} = value of feature f_1 for data element d_1

f_{12} = value of feature f_1 for data element d_2

f_{21} = value of feature f_2 for data element d_1

f_{22} = value of feature f_2 for data element d_2

Algorithm Steps

For a given test instance x :

1. Calculate the distance between x and all points in the training set.
2. Sort the training points by distance from x .
3. Select the k nearest neighbors.
4. For classification, perform a majority vote among the k neighbors. For regression, average the k neighbors' values.

Choosing the Right k

Selecting an appropriate value for k is crucial:

Small k : Can be noisy and lead to overfitting.

Large k : May smooth out the predictions too much and underfit the data.

A common approach is to use cross-validation to determine the optimal k value.

Summary

k -Nearest Neighbors is a straightforward and versatile algorithm useful for both classification and regression. Its ease of implementation and lack of training phase make it an attractive choice for many problems, though its performance can be impacted by the size and dimensionality of the dataset. When using k -NN, careful consideration must be given to the choice of k and the distance metric to ensure optimal results.

diabetes.csv

```
In [1]: #import Libraries
import pandas as pd
```

```
from sklearn.model_selection import train_test_split
from sklearn.neighbors import KNeighborsClassifier
```

```
In [2]: data = pd.read_csv("diabetes.csv")
data
```

```
Out[2]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age
0	6	148	72	35	0	33.6	0.627	50
1	1	85	66	29	0	26.6	0.351	31
2	8	183	64	0	0	23.3	0.672	32
3	1	89	66	23	94	28.1	0.167	21
4	0	137	40	35	168	43.1	2.288	33
...
763	10	101	76	48	180	32.9	0.171	63
764	2	122	70	27	0	36.8	0.340	27
765	5	121	72	23	112	26.2	0.245	30
766	1	126	60	0	0	30.1	0.349	47
767	1	93	70	31	0	30.4	0.315	23

768 rows × 9 columns



```
In [3]: data.shape
```

```
Out[3]: (768, 9)
```

```
In [4]: data.isna().sum()
```

```
Out[4]: Pregnancies      0
Glucose      0
BloodPressure  0
SkinThickness  0
Insulin      0
BMI          0
DiabetesPedigreeFunction  0
Age          0
Outcome      0
dtype: int64
```

```
In [5]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 768 entries, 0 to 767
Data columns (total 9 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Pregnancies           768 non-null   int64
1   Glucose               768 non-null   int64
2   BloodPressure         768 non-null   int64
3   SkinThickness         768 non-null   int64
4   Insulin               768 non-null   int64
5   BMI                   768 non-null   float64
```

```

6 DiabetesPedigreeFunction 768 non-null float64
7 Age 768 non-null int64
8 Outcome 768 non-null int64
dtypes: float64(2), int64(7)
memory usage: 54.1 KB

```

```

In [6]: #Segregating predictor variables
x = data.iloc[:, 0:8]
x

```

```

Out[6]:
   Pregnancies  Glucose  BloodPressure  SkinThickness  Insulin  BMI  DiabetesPedigreeFunction  Age
0            6      148             72             35         0  33.6                0.627      50
1            1       85             66             29         0  26.6                0.351      31
2            8      183             64              0         0  23.3                0.672      32
3            1       89             66             23        94  28.1                0.167      21
4            0      137             40             35       168  43.1                2.288      33
...         ...      ...             ...             ...         ...      ...                ...      ...
763          10      101             76             48       180  32.9                0.171      63
764           2      122             70             27         0  36.8                0.340      27
765           5      121             72             23       112  26.2                0.245      30
766           1      126             60              0         0  30.1                0.349      47
767           1       93             70             31         0  30.4                0.315      23

```

768 rows × 8 columns

```

In [7]: #Segregating the target/class variable
y = data['Outcome']
y

```

```

Out[7]:
0      1
1      0
2      1
3      0
4      1
..
763    0
764    0
765    0
766    1
767    0
Name: Outcome, Length: 768, dtype: int64

```

```

In [8]: #split into training and test datasets
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size = 0.2, random_state=

```

```

In [9]: print(x_train.shape)
print(x_test.shape)
print(y_train.shape)
print(y_test.shape)

```

```

(614, 8)
(154, 8)

```

$$\begin{pmatrix} 614, \\ 154, \end{pmatrix}$$

```
In [22]: #kNN Classifier with k=27 means 27 closest neighbours are considered.  
nn = KNeighborsClassifier(n_neighbors=15)
```

```
In [23]: #Train the classifier with the training data
         model = nn.fit(x_train, y_train)
```

```
In [24]: prediction = model.predict(x_test)
```

```
In [25]: prediction
```

```
Out[25]: array([1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 0, 1, 0, 0, 0, 0, 0, 0,
                1, 1, 0, 1, 1, 1, 0, 0, 0, 0, 0, 1, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0,
                0, 0, 1, 1, 0, 0, 0, 0, 1, 0, 0, 0, 1, 0, 0, 0, 1, 0, 1, 0, 0, 0,
                0, 1, 0, 1, 1, 1, 1, 0, 1, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0, 1, 1, 0,
                1, 0, 0, 0, 0, 0, 1, 0, 0, 0, 1, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0,
                1, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 1, 1, 1, 0,
                1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 1, 0, 0, 0, 0, 0, 1, 0, 0],
                dtype=int64)
```

```
In [26]: diff=pd.DataFrame({"Actual":y_test,"Prediction":prediction})
diff
```

Out[26]:

Actual	Prediction
--------	------------

285	0	1
101	0	0
581	0	0
352	0	0
726	0	0
...
563	0	0
318	0	0
154	1	1
684	0	0
643	0	0

154 rows × 2 columns

```
In [17]: #To Store the data of above dataframe to a csv file
diff.to_csv('diabetes_data.csv')
```

```
In [16]: x_test
```

Out[16]:	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age
285	7	136	74	26	135	26.0	0.647	51
101	1	151	60	0	0	26.1	0.179	22

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age
581	6	109	60	27	0	25.0	0.206	27
352	3	61	82	28	0	34.4	0.243	46
726	1	116	78	29	180	36.1	0.496	25
...
563	6	99	60	19	54	26.9	0.497	32
318	3	115	66	39	140	38.1	0.150	28
154	8	188	78	0	0	47.9	0.137	43
684	5	136	82	0	0	0.0	0.640	69
643	4	90	0	0	0	28.0	0.610	31

154 rows × 8 columns

Confusion Matrix

- There are four possibilities with regards to the cricket match win/loss prediction:
 - 1. the model predicted win and the team won (TP)
 - 2. the model predicted win and the team lost (FP)
 - 3. the model predicted loss and the team won (FN)
 - 4. the model predicted loss and the team lost (TN)

```
In [27]: #Metric Confusion Matrix
from sklearn.metrics import confusion_matrix
cm = confusion_matrix(y_test, prediction)
cm
```

```
Out[27]: array([[89, 10],
               [23, 32]], dtype=int64)
```

```
In [146... TN = cm[0][0]
FP = cm[0][1]
FN = cm[1][0]
TP = cm[1][1]
print(TP, FN, TN, FP)
```

28 27 88 11

Model Accuracy

$$\text{Model accuracy} = \frac{TP + TN}{TP + FP + FN + TN}$$

	ACTUAL WIN	ACTUAL LOSS
Predicted Win	85	4
Predicted Loss	2	9

In context of the above confusion matrix, total count of TPs = 85, count of FPs = 4, count of FNs = 2 and count of TNs = 9.

$$\therefore \text{Model accuracy} = \frac{TP + TN}{TP + FP + FN + TN} = \frac{85 + 9}{85 + 4 + 2 + 9} = \frac{94}{100} = 94\%$$

Model Accuracy/Accuracy Score

```
In [147... Model_Accuracy=(TP+TN)/(TP+TN+FN+FP)
print("Accuracy Score:",Model_Accuracy)
```

Accuracy Score: 0.7532467532467533

```
In [148... from sklearn.metrics import accuracy_score
Accuracy=accuracy_score(y_test,prediction)
Accuracy
```

Out[148... 0.7532467532467533

Error Rate

$$\text{Error rate} = \frac{\text{FP} + \text{FN}}{\text{TP} + \text{FP} + \text{FN} + \text{TN}} = \frac{4 + 2}{85 + 4 + 2 + 9} = \frac{6}{100} = 6\%$$

= 1 - Model accuracy

Error Rate

In [149...

```
Error_Rate=1-Model_Accuracy  
print("Error Rate:", Error_Rate)
```

Error Rate: 0.24675324675324672

Sensitivity

- The sensitivity of a model measures the proportion of TP examples or positive cases which were correctly classified.

$$\text{Sensitivity} = \frac{\text{TP}}{\text{TP} + \text{FN}} = \frac{85}{85 + 2} = \frac{85}{87} = 97.7\%$$

Sensitivity

In [150...

```
Sensitivity= TP / (TP + FN)  
print("Sensitivity:",Sensitivity)
```

Sensitivity: 0.509090909090909

Specificity

- Specificity of a model measures the proportion of negative examples which have been correctly classified.

$$\text{Specificity} = \frac{\text{TN}}{\text{TN} + \text{FP}} = \frac{9}{9 + 4} = \frac{9}{13} = 69.2\%$$

Specificity


```
In [151... Specificity= TN / (TN + FP)
print("Specificity:",Specificity)
```

Specificity: 0.8888888888888888

To find the best value of k for highest accuracy_score.

```
In [28]: k=[]
for i in range (1,30):
    nn = KNeighborsClassifier(n_neighbors=i)
    model = nn.fit(x_train, y_train)
    prediction = model.predict(x_test)
    from sklearn.metrics import accuracy_score
    k.append(accuracy_score(y_test,prediction))

import matplotlib.pyplot as plt
plt.plot(range(1,30),k,marker="*")
plt.show()
print(k)
```



```
[0.7207792207792207, 0.7272727272727273, 0.7402597402597403, 0.7207792207792207, 0.7337662337662337, 0.7727272727272727, 0.7597402597402597, 0.7662337662337663, 0.7662337662337663, 0.7597402597402597, 0.7467532467532467, 0.7727272727272727, 0.7662337662337663, 0.7727272727272727, 0.7857142857142857, 0.7857142857142857, 0.7857142857142857, 0.7727272727272727, 0.7727272727272727, 0.7727272727272727, 0.7662337662337663, 0.7727272727272727, 0.7662337662337663, 0.7792207792207793, 0.7532467532467533, 0.7727272727272727, 0.7532467533, 0.7662337662337663, 0.7662337662337663]
```

tshirt.csv

```
In [153... #import Libraries
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.neighbors import KNeighborsClassifier
```

```
In [154... df=pd.read_csv('tshirt.csv')
df
```

```
Out[154...      Height  Weight  Size
0      158      58    M
1      158      59    M
2      158      63    M
3      160      59    M
4      160      60    M
5      163      60    M
6      163      61    M
7      160      64    L
8      163      64    L
9      165      61    L
10     165      62    L
11     165      65    L
12     168      62    L
13     168      63    L
14     168      66    L
15     170      63    L
16     170      64    L
17     170      68    L
```

```
In [155... #Segregating predictor variables
x = df.iloc[:, 0:2]
y = df.iloc[:,2]
```

```
In [156... #Alternate Method to take x & y (Segregating predictor variables)
x=df[['Height','Weight']]
y=df['Size']
```

```
In [157... x.shape
```

```
Out[157... (18, 2)
```

In [158... `y.shape`

Out[158... `(18,)`

In [159... `#split into training and test datasets`
`# x_train, x_test, y_train, y_test = train_test_split(x, y, test_size = 0.2, random_state=42)`

In [160... `# print(x_train.shape)`
`# print(x_test.shape)`
`# # print(y_train.shape)`
`# print(y_test.shape)`

In [161... `#kNN Classifier with k=9 means 9 closest neighbours are considered`
`nn = KNeighborsClassifier(n_neighbors=9)`

In [162... `#Train the classifier with the training data`
`model = nn.fit(x, y)`

In [163... `prediction = model.predict(x)`

In [164... `prediction`

Out[164... `array(['M', 'M', 'M', 'M', 'M', 'M', 'L', 'M', 'L', 'L', 'L', 'L', 'L', 'L', 'L', 'L', 'L', 'L'], dtype=object)`

In [165... `diff=pd.DataFrame({'Actual':y, "Predicted":prediction})`
`diff`

Out[165...

	Actual	Predicted
0	M	M
1	M	M
2	M	M
3	M	M
4	M	M
5	M	M
6	M	L
7	L	M
8	L	L
9	L	L
10	L	L
11	L	L
12	L	L
13	L	L
14	L	L
15	L	L

	Actual	Predicted
0	M	M
1	M	M
2	M	M
3	M	M
4	M	M
5	M	M
6	M	L
7	L	M
8	L	L
9	L	L
10	L	L
11	L	L
12	L	L
13	L	L
14	L	L
15	L	L

	Actual	Predicted
16	L	L
17	L	L

```
In [166... #Metric Confusion Matrix
from sklearn.metrics import confusion_matrix
cm = confusion_matrix(y, prediction)
cm
```

```
Out[166... array([[10,  1],
        [ 1,  6]], dtype=int64)
```

```
In [167... TN=cm[0][0]
TP=cm[1][1]
FN=cm[1][0]
FP=cm[0][1]

print(TN,TP,FN,FP)

10 6 1 1
```

```
In [78]: Model_Accuracy=(TP+TN)/(TP+TN+FN+FP)
print("Accuracy:",Model_Accuracy)

Accuracy: 0.8888888888888888
```

```
In [168... from sklearn.metrics import accuracy_score
print("Accuracy Score: ",accuracy_score(y,prediction))

Accuracy Score: 0.8888888888888888
```

```
In [79]: Error_Rate=1-Model_Accuracy
print("Error Rate:", Error_Rate)

Error Rate: 0.11111111111111116
```

```
In [80]: Sensitivity= TP / (TP + FN)
print("Sensitivity:",Sensitivity)

Sensitivity: 0.8571428571428571
```

```
In [81]: Specificity= TN / (TN + FP)
print("Specificity:",Specificity)


Specificity: 0.9090909090909091
```

To find the best value of k for highest accuracy_score.

```
In [169... k=[]
for i in range (1,18):
    nn = KNeighborsClassifier(n_neighbors=i)
    model = nn.fit(x,y)
    prediction = model.predict(x)
    from sklearn.metrics import accuracy_score
    k.append(accuracy_score(y,prediction))

import matplotlib.pyplot as plt
plt.plot(range(1,18),k,marker="*")
```

```
plt.show()
print(k)
```



```
[1.0, 0.9444444444444444, 1.0, 0.8888888888888888, 0.8888888888888888, 0.8888888888888888, 0.8888888888888888, 0.9444444444444444, 0.8333333333333334, 0.8888888888888888, 0.8333333333333334, 0.8888888888888888, 0.8333333333333334, 0.8888888888888888, 0.6111111111111112, 0.6111111111111112, 0.6111111111111112]
```

AptitudeCommunication.csv

In [171...

```
import pandas as pd
df=pd.read_csv('AptitudeCommunication.csv')
df
```

Out[171...

	Name	Aptitude	Communication	Class
0	Karuna	2	5.0	Speaker
1	Bhavan	2	6.0	Speaker
2	Gaurav	7	6.0	Leader
3	Parul	7	2.5	Intel
4	Dinesh	8	6.0	Leader
5	Jani	4	7.0	Speaker
6	Bobby	5	3.0	Intel
7	Parimal	3	5.5	Speaker
8	Govind	8	3.0	Intel
9	Sushant	6	5.5	Leader
10	Gauri	6	4.0	Intel
11	Bharat	6	7.0	Leader
12	Rajvi	6	2.0	Intel
13	Pradip	9	7.0	Leader

In [172...

```
x=df[['Aptitude','Communication']]
y=df['Class']
```

In [173... `x.shape`

Out[173... `(14, 2)`

In [174... `y.shape`

Out[174... `(14,)`

In [180... `from sklearn.neighbors import KNeighborsClassifier`

```
nn=KNeighborsClassifier(n_neighbors=5)
model=nn.fit(x,y)
y_pred=model.predict(x)
y_pred
```

Out[180... `array(['Speaker', 'Speaker', 'Leader', 'Intel', 'Leader', 'Speaker',
 'Intel', 'Speaker', 'Intel', 'Leader', 'Intel', 'Leader', 'Intel',
 'Leader'], dtype=object)`

In [181... `diff=pd.DataFrame({"Actual":y,"Predicted":y_pred})`
`diff`

Out[181...

	Actual	Predicted
--	--------	-----------

0	Speaker	Speaker
1	Speaker	Speaker
2	Leader	Leader
3	Intel	Intel
4	Leader	Leader
5	Speaker	Speaker
6	Intel	Intel
7	Speaker	Speaker
8	Intel	Intel
9	Leader	Leader
10	Intel	Intel
11	Leader	Leader
12	Intel	Intel
13	Leader	Leader

In [182... `#Prediction for specific values of Aptitude & Communication`
`prediction=model.predict([[5,4.5]])`
`prediction`

Out[182... `array(['Intel'], dtype=object)`

In [183... `#Confusion Matrix`
`from sklearn.metrics import confusion_matrix`


```
cm = confusion_matrix(y, y_pred)
cm
```

```
Out[183...] array([[5, 0, 0],
        [0, 5, 0],
        [0, 0, 4]], dtype=int64)
```


```
In [188...] from sklearn.metrics import accuracy_score
print("accuracy_score:", accuracy_score(y,y_pred))
```

```
accuracy_score: 1.0
```

To find the best value of k for highest accuracy_score

```
In [190...] k=[]
for i in range (1,15):
    nn = KNeighborsClassifier(n_neighbors=i)
    model = nn.fit(x,y)
    prediction = model.predict(x)
    from sklearn.metrics import accuracy_score
    k.append(accuracy_score(y,prediction))

import matplotlib.pyplot as plt
plt.plot(range(1,15),k,marker="*")
plt.show()
print(k)
```



```
[1.0, 1.0, 1.0, 1.0, 1.0, 1.0, 1.0, 1.0, 1.0, 1.0, 0.8571428571428571, 0.5714285714285714, 0.42857142857142855, 0.42857142857142855, 0.35714285714285715]
```

```
In [ ]:
```

```
In [ ]:
```